

FirstRound AI — Interview Scorecard

Recruiter artifact. Not shown to the candidate-facing interview UI.

Candidate: Muhammad Taha Farrukh

Role: Junior AI Engineer

Interview date: 2026-08-14

Duration: 482 seconds

Status: partial

Sample mode:

Decision

Overall score (1–5, evidence-gated): 3.6

Mean confidence: 0.91

Recommendation: borderline

Reasoning: Mixed or borderline evidence; not enough to recommend GO or NO_GO.

Legacy internal score (0–100): 62 (REVIEW)

Competencies

Competency	Score	Confidence	Evidence quote	Reasoning
Technical competence	4	0.91	I compared chunking strategies based on retrieve quality and context preservation recursive corrector spreading is a strong general purpose base line because it tries to preserve boundaries.	Mixed or borderline evidence; not enough to recommend GO or NO_GO.
Problem solving	2	0.91	I would raise that pipeline rather than immediately changing the prompt first I would inspect the similarity scores and verify chunk boundaries and embaring then inspect the exact context passed to the LM.	Mixed or borderline evidence; not enough to recommend GO or NO_GO.
Communication	4	0.91	The male limitation is that each session can consume memory and the data is in suitable as a shared persistent production store. For production I would separate injection from queuering persistent bearings in a proper vector store reuse the embedding model and introduce monitoring and cashing we are appropriate.	Mixed or borderline evidence; not enough to recommend GO or NO_GO.
Project understanding	4	0.91	The male limitation is that each session can consume memory and the data is in suitable as a shared persistent production store. For production I would separate injection from queuering persistent bearings in a proper vector store reuse the embedding model and introduce monitoring and cashing we are appropriate.	Mixed or borderline evidence; not enough to recommend GO or NO_GO.
GitHub credibility	4	0.91	The male limitation is that each session can consume memory and the data is in suitable as a shared persistent production store. For production I would separate injection from queuering persistent bearings in a proper vector store reuse the embedding model and introduce monitoring and cashing we are appropriate.	Mixed or borderline evidence; not enough to recommend GO or NO_GO.

JD / resume fit	—		None	Not scored: the transcript does not contain a usable candidate quote for this competency.
-----------------	---	--	------	---

Strengths

- Phase 2 gap analysis shows strong skill matches with the JD.

Concerns

- q4: shallow or under-specified answers.
- q5: shallow or under-specified answers.
- q6: shallow or under-specified answers.
- q9: off-topic response.
- JD skills are missing from the approved profile.
- Some JD skills are only weakly evidenced on the resume.

Guardrail flags

- insufficient_evidence: No valid candidate-transcript quote of at least 12 characters.
- partial_interview: Only 9 of 12 approved questions were attempted.

GitHub grounding

GitHub-grounded questions asked: 3
Questions attempted: 9 of 12

Limitations

Competency scores are refused unless a candidate-transcript quote of at least 12 characters appears verbatim. This interview may be partial. LiveKit reconnection after a dropped call is not implemented; SQLite restores controller state only. This PDF is a recruiter artifact and is not shown on the candidate frontend.